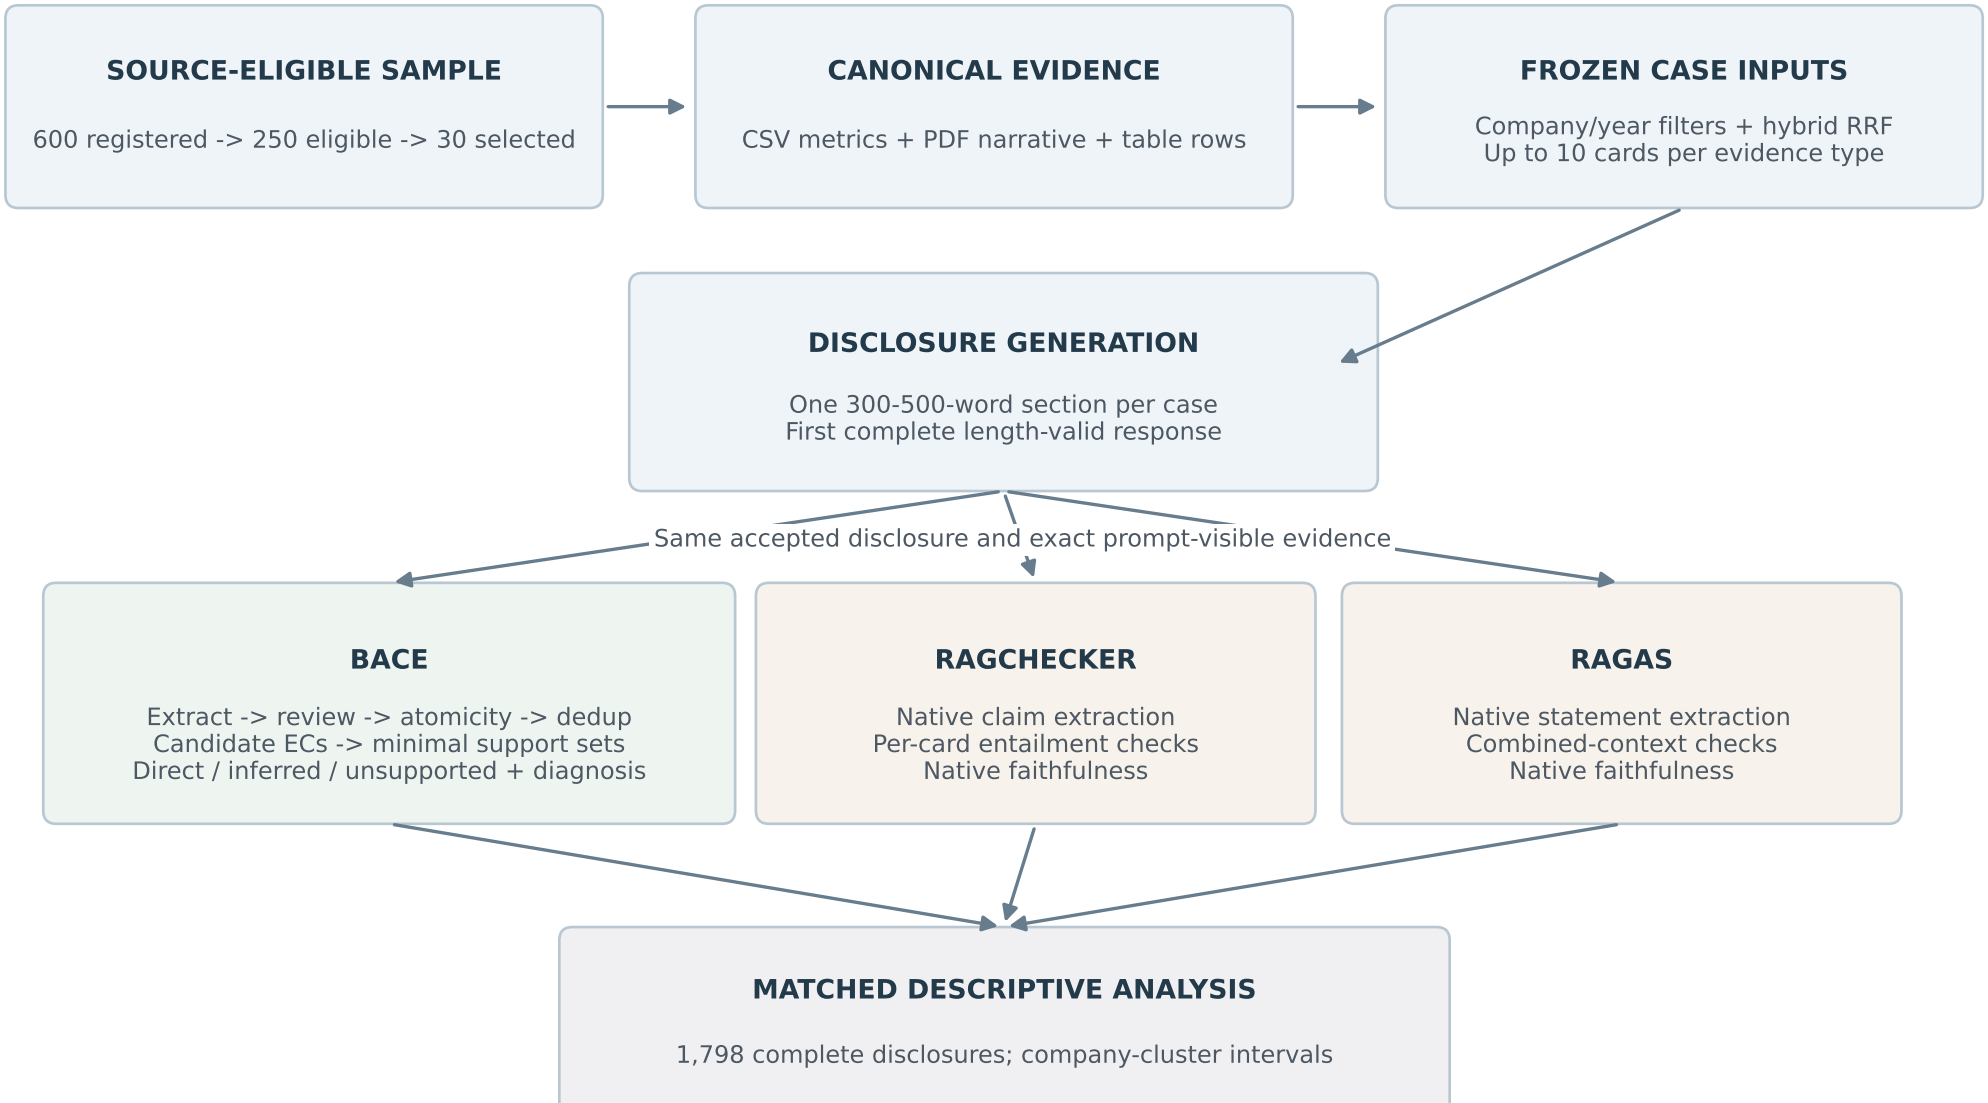
SOURCE-ELIGIBLE SAMPLE
600 registered -> 250 eligible -> 30 selected
CANONICAL EVIDENCE
CSV metrics + PDF narrative + table rows
FROZEN CASE INPUTS
Company/year filters + hybrid RRF
Up to 10 cards per evidence type
DISCLOSURE GENERATION
One 300-500-word section per case
First complete length-valid response
Same accepted disclosure and exact prompt-visible evidence
BACE
Extract -> review -> atomicity -> dedup
Candidate ECs -> minimal support sets
Direct / inferred / unsupported + diagnosis
RAGCHECKER
Native claim extraction
Per-card entailment checks
Native faithfulness
RAGAS
Native statement extraction
Combined-context checks
Native faithfulness
MATCHED DESCRIPTIVE ANALYSIS
1,798 complete disclosures; company-cluster intervals